The ESL Programming Language

Reference Manual

Brian G. Lucas

Version of 2012-10-11

Table of Contents

1 Introduction

ESL is a new programming language designed to be used for efficient programming of embedded and other "small" systems. ESL is an acronym for Embedded Systems Language (pronounced: "ESS-el").

ESL is a typed compiled language with features that allow the programmer to dictate the concrete representation of data values. This distinguishes it from languages which implement only "abstract" types or types whose representation is architecture-dependent. The programmer can dictate the details of data representation, including such things as "endian-ness" and the exact placement of bits, which are necessary in dealing with external representations of data layout, e.g., communication protocols or device registers.

ESL is not really a "new" programming language - all the elements have probably been seen in other programming languages. In many respects, it is a conventional language whose features, for the most part, will be familiar to programmers who have had experience with any of the compiled languages introduced in the past 50 years.

If the ESL syntax bears some resemblance to Google's new programming language *Go*, that is due to common inspiration. The syntax for ESL was mostly in place before G*o* was made public. There is one exception: the syntax for methods. This was shamelessly stolen from *Go* and incorporated into ESL.

2 Syntactic Elements

2.1 Comments

There are two ways to comment text: by line or by block. Line comments start with "`//`" and end at the end of the line. Block comments start with "`/*`" and end with "`*/`". Block comments do not nest.

2.2 Reserved Words

There are <u>no reserved words</u> in ESL. There are many keywords that have special meaning only in

specific contexts. Therefore, these keywords can be used as ordinary identifiers. Keywords are always lower case.

2.3 Identifiers

The first character of an identifier must be from the set {A-Z,a-z}, remaining characters can be chosen from the set {A-Z,a-z,0-9,_}. The system reserves identifiers starting with a "_"for predefined symbols. In addition, the identifier consisting of a single "_" has a special meaning. It is not entered into the symbol table and can be used to specify anonymous fields in records, or placeholders in enumerations.

Identifiers may also contain UTF8 sequences any place an alphabetic character is valid. (Identifiers with UTF8 sequences may not be handled correctly in the LLVM backend, by assemblers, or by linkers.)

Identifiers may be keywords depending on context. Otherwise they are, as yet, undeclared names or previously declared names of a several kinds: type names, variable names, procedure names, package names, field names, enumeration constant names, or aliases.

2.3.1 Identifier Namespaces

There are two visibility classes of identifiers: public and private. The public identifiers are nested at four levels: universal, global, package, and procedure. Universal identifiers are pre-defined by the compiler. Global identifiers declared at the outermost nesting level of a program. Package identifiers are declared within a package. Packages may be nested. Finally, procedure identifiers are declared within a procedure. There is no further nesting within a procedure, i.e., no "blocks", and procedures can not be nested.

Private namespaces include enumeration constants, fields within a record, and methods. They are accessed by prefixing the enumeration type name, the record variable name, or a variable of the method type.

2.4 Numeric Literals

Numbers are by default decimal. For other number bases, a two-character prefix is used. The first character is always a zero, the second character, always lower case, indicates the number base:

    0b binary - Only digits 0-1 are allowed.
    0o octal - Only digits 0-7 are allowed.
    0x hexadecimal - Only digits 0-9 and letters A-F or a-f are allowed.
In any number, after the first digit (which follows the optional prefix), an underline "_" is allowed for readability.

Examples:
```
399                   // decimal
0b0010_0001           // binary
0o02_56               // octal
0x01_F4               // hexadecimal
1_999_999             // decimal
```

2.5 Character Literals

Character literals are unsigned integer constants. Character literals are enclosed within single quotes, the literal is either any single ASCII character (excluding the single quote and newline) or an escape sequence. Escape sequences start with a backslash and include:

    \\ - represents a single backslash
    \n - represents a newline
    \r - represents a carriage return
    \f - represents a form feed
    \t - represents a horizontal tab
    \b - represents a backspace
    \v - represents a vertical tab
    \x*XX* - represents an 8-bit character with value given by the two hex digits *XX*
    \u*XXXX* - represents a 16-bit (e.g. unicode) character with value given by four hex digits
    \U*XXXXXXXX* - represents a 32-bit (e.g. unicode) character with value given by eight hex digits

2.6 String Literals

A string literal is a constant array of bytes terminated by a NUL. Each byte is either an ASCII character, an 8-bit escape as defined in the above character literal section, or a part of a UTF8 sequence. The character escapes which generate greater than 8-bit values are converted into a sequence of bytes by UTF8 encoding.

3 Types and Declarations

Types in ESL are not abstract in the sense that they are designed to be concrete descriptions of containers for a set of values. When instances of types are allocated, the assignment to bit positions and memory addresses depends on architecture-dependent defaults and which can be overridden by the use of optional type attributes.

There are two units of allocation:

- *bits* - the atomic unit of data
- *bytes* - the smallest sequence of bits that has a memory address, in most contemporary machines this is 8 bits.

Types consist of the integral types, references (pointers), and aggregate types. Integral types are unsigned or signed and are always sub-ranges. Enumerations are a special type of unsigned type. Signed integers are assumed to have twos complement representation. Aggregate types are either arrays or records.

Enumerations and records can be "extended" from base types.   Details on how each of these types are extended will be discussed in the section where the specific type is defined.

Types are declared with type statements. Type statements may occur anywhere in the statement flow.

3.1 Integer Type

The integer types are specified by giving their lowest and highest values. The resulting type may be signed or unsigned, depending on the values given. The representation of unsigned sub-ranges always includes zero. The representation of signed sub-ranges is symmetrical (in the twos-complement sense)

around zero.

Examples:

```
type percent: 0..100;          // an unsigned type
type srange: -1000..1000;      // a signed type
```

3.2 Enumerated Type

Enumerated types are a mapping onto the unsigned integers. They are defined by listing all the values, in monotonic increasing order, that make up the type.

These values may represented by identifiers. The identifiers will be assigned specific numeric values by the compiler, or optionally, can be given specific numeric values by the programmer. An enumerated type consists of a contiguous range of values from zero to some maximum value, there are no gaps. All of the values need not be represented by identifiers.

Operations on enumerated types are limited to assignment and comparison. However, casting to unsigned integer types is allowed.

Each enumerated type is distinct. That is, two different enumerated types are not compatible for assignment or comparison.

An enumerated type can be extended by adding additional values.

Examples:

```
type answer: (yes, no);                    // yes=0, no=1
type foo: (low=5, medium, high=10)// medium=6,
          // values 0..4, 7..9 are part of the enumeration
type fuzzy(answer): (maybe, possibly);
          // extends answer: yes=0, no=1, maybe=2, possibly=3
```
The size of an enumeration is the minimal number of bytes necessary to hold the largest defined value. For all the enumerated types in the previous example, the size that would be needed for all three, i.e., answer, foo, and fuzzy, would be one byte in most architectures.

3.3 Data Reference Type

This is just a pointer to another type, the base type. The base type must have a type name.

It is possible that a programmer wants to reference a type that has not yet been declared. This is possible by referencing the type name and then later later completing the type definition for that name.

Examples:

```
type byte: 0..255;              // the base type
type ptr_to_byte: @byte;        // a pointer to the base type
type ptr_to_forward: @forward;  // forward reference
type forward: sometype;         // the real declaration is later
```

3.4 Procedure Reference Type

See section 7.5.

3.5 Array Type

Arrays are aggregates of a base type. The base type is accessed by the indexing operation. Array indices are unsigned and always start at zero.

Array slicing is an operation that allows access to a contiguous sub-array of values. A slice operator is square brackets enclosing an offset and length separated by a colon.

3.5.1 Fixed Sized Arrays

The index specification of a fixed sized array can be either an expression that is a compile-time constant, or the name of an unsigned type (e.g. an enumeration).

Examples:

```
type array1: [100]_uint;
type enum: (zero, one, two, three);
type array2: [enum]enum;
```

3.5.2 Unknown Sized Arrays

Arrays of indeterminate size are allowed as targets of a reference (pointer).

Examples:

```
type pstring: @[]_byte;
type argv: @[]@[]_byte;              // Unix argv type
```

3.6 Record Type

Record types are a way to define an aggregation of dissimilar types as new type. A record is defined as an order list of *fields*. Each field may be of any type. The following example demonstrates how the type base is declared with two fields: one and two.

```
type base:
{   one: _uint32;
    two: _uint32;
};
```

The final field of a record may have a type of unknown size:

```
type ident:
{   length: _uint;
    name:   []_byte;
};
```

As mentioned earlier, a record type is one of the types that can be "extended". The following example

extends the type 'base' declared above:.

```
type newbase(base):
{   three: _uint16;
    four:  _uint8;
};
```

The new extended type, 'newbase', includes all the fields that were declared in the type 'base', and two additional fields. New fields in the extended type are always added at the end.

3.7 Type Attributes

Type attributes can be used to alter the way instances of that type are allocated or behave. Allocation attributes alter the size, alignment, bit-order, and byte-order of the memory structure. Usage attributes alter the way that instances of the type can be used.

3.7.1 Size

A size attribute overrides the default (target dependent) size of a type:

bits(*cexpr*) - instances of this type should be exactly this many bits in size.

size(*cexpr*) - similar to bits(), except the size is in memory-units.

At most, only one of these attributes should be present. If no alignment attributes are present, the size of a type may be increased to a target-dependent value.

3.7.2 Alignment

Alignment attributes change the way a type is positioned in memory:

packed - when the type is allocated no padding occurs at the bit level nor the memory-unit level.

nopad - that padding occurs only up to the next memory-unit.

align(*cexpr*) - padding up to an alignment in memory, or a record offset, will be made to the *cexpr* byte boundary.

At most, only one of these attributes should be present. If none are present, then allocation will use whatever the target architecture prefers. On many architectures alignment is type dependent .

3.7.3 Bit Order

The bit order attributes affect allocation within an aggregate, record or array, that is packed.

lsb - allocation begins at the least significant bit within a memory-unit and with the memory-unit with the lowest address, i.e. "little-endian"

msb - allocation begins at most significant bit within a memory-unit and with the memory-unit with

the highest address, i.e. "big-endian" also known as "network order"

At most, only one of these attributes should be present. If neither attribute is given, the compiler will use whatever the target architecture prefers.

3.7.4 Memory Order

The memory order attributes affect how a multiple byte type is stored in memory. The compiler will "byte-swap", if necessary when loading or storing to insure that the byte-order is correct for the target machine.

`le` - types that span multiple bytes will have the least significant byte at the first allocated byte.

`be` - types that span multiple bytes will have the most significant byte at the first allocated byte.

At most, only one of these attributes should be present. If neither attribute is given, the compiler will use whatever the target architecture prefers.

3.7.5 Access Restrictions

These attributes allow the programmer to restrict access to allocations of this type.

`ro` - read only

`wo` - write only

At most, only one of these attributes should be present. If neither attribute is given, the default is to allow both read and write.

3.7.6 Input and Output

These attributes indicate to the compiler that allocations of this type has external side effects. These attributes are most often used to describe device registers which do not behave as normal memory.

`in` - loads of instances of this type must not be optimized away because reading may cause external side effects

`out` - stores to instances of this type must not be optimized away because writing may cause external side effects

Some C-based programming languages use the keyword "volatile" to lump together both attributes `in` and `out`.

3.8 Built-in Types

There are a small number of pre-defined types. These include:

`_boolean` is an enumerated type consisting of the constants `false` and `true`.

_byte is an unsigned integral type with a range equal to that which can be stored in the target architecture's minimal addressable unit. For modern processors, this is equivalent to _uint8.

_uint8,_uint16,_uint32,_uint64 are unsigned types of 8, 16, 32, and 64 bits respectively.

_int8,_int16,_int32,_int64 are signed types of 8, 16, 32, 64 bits respectively.

_uint and _int are unsigned and signed integral types with a range equal to that which can be stored in the target architecture's "natural word size". This is usually the size of a integer register.

_uintptr is an unsigned integral type the same size as an "address" (i.e., "pointer") in the target architecture.

_memory is a synonym for the []_byte type.

_address is a synonym for the @_memory type.

Because these type identifiers start with "_" they can not be redefined by the programmer. For convenience, the alias "boolean" is predefined to be the same as "_boolean"; but can be redefined by the programmer. The original boolean constants are always available by the names _boolean.false and _boolean.true.

4   Variables

4.1 Variable Declarations

Variable declarations instantiate a type. The visibility and lifetime of the variable depend on the nesting level at which the declaration appears.

4.2 Variable Attributes

4.2.1 External

The external attribute indicates that the variable is allocated external to the program being compiled. The external name (i.e. linkage name) is the variable name (without any package prefixes) unless specified otherwise. External variables can also be placed at a specific address.

Examples:
```
var foo: _int: external;                    // linkage name is "foo";
var bar: _int: external("foobar");          // linkage name is "foobar"
var vectors: [32]_uint: external(0x20000000); // at a specific address
```

4.2.2 Global

The `global` attribute indicates that the current allocation of a variable should be available, by name, to programs external to the program being compiled. The linkage name is the variable name (without any package prefixes) unless specified otherwise.

Examples:

```
var foo: _int: global;                  // linkage name is "foo";
var bar: _int: global("foobar");        // linkage name is "foobar"
```

4.2.3 Segment

[Not yet implemented.]

5   Expressions

5.1 Type Inference

In many cases, as an expression is being evaluated the type of an identifier can be inferred. This inference takes place in the following situations:

- Evaluation of the right-hand side of an assignment

- Evaluation of actual parameter in a procedure call

- Evaluation of a returned value in a procedure

As a simple example of the first case, take the assignment of an enumeration constant to an enumeration variable:

```
type enum: (A, B, C);
var foo: enum;
foo = A;                 // same as foo = enum.A
```

Since enumeration constants are in a private namespace associated with the enumeration type, the name "A" would be invisible, it would have to be referred to by its qualified name "enum.A". However, assignment allows the type of the left-hand side (the "target type") to guide the type of the right-hand side. Any names on the right hand side which are unknown are interpreted in the context of the type.

The target type is also set when evaluating an expression for an actual parameter or returned value. Given the above definition of "enum", and example of the second case is:

```
proc bar(a: enum);
proc baz()
{   bar(B);              // same as bar(enum.B)
}
```

If the target type has an aggregate type, type inference controls the construction of aggregate expressions.

5.2 Type Conversions

Inevitably, one will have to convert the type of an expression into one of more preferable type for a particular situation. In some cases the conversion will be done implicitly by the compiler. In other cases the programmer must explicitly cause a type conversion.

5.2.1 Explicit

5.2.1.1 Ordinary Casts

Ordinary casts convert the type of a value to a different type. The name of the desired type is used. Some examples:

```
type Enum: (A, B, C, D, E, F);
var ev: Enum;
var x: _uint8;
x = _uint8(ev);        // convert from enumeration to unsigned int
ev = Enum(x+1);        // convert from unsigned int to enumeration

type pFoo: @Foo;
type pBar: @Bar;
var pf: pFoo;
var pb: pBar;
pf = pFoo(pb);         // convert from @Bar to @Foo
```

5.2.1.2 Container Casts

Container casts are used to convert from references (pointers) to a type to a reference to an aggregate that contains that type. This is best demonstrated by an example:

```
type Foo:
{   count: _uint;
    bar:   Bar;        // some previously declared type Bar
};
type pFoo: @Foo;
. . .
var p: @Bar;           // a pointer to Foo.bar
var pf: pFoo;          // a pointer to the container, Foo
. . .
pf = pFoo(p, .bar);            // convert to pointer to container
```

A somewhat more complicated example follows:

```
type Data:
{   len: _uint;
    buf: []_byte;
};
type pData: @Data;
. . .
var n: _uint;          // an index into Data.buf[]
var p: @_byte;         // a pointer to Data.buf[n]
var pd: pData;         // a pointer to the container, Data
. . .
pd = pData(p, .buf[n]);          // convert to pointer to container
```

5.2.2 Implicit

Implicit type conversions occur at "assignment points". An assignment point occurs:

- when assigning a value as a result of an assignment statement

- when assigning a value to a parameter in a procedure call

- when assigning a value in a `return` statement

Implicit conversions are of two differing kinds: width adjustment and reference coercion.

5.2.2.1 Widening versus Truncation

Probably the most common situation is the conversion of a variable from one bit width size to another.
Widening of scalars occurs automatically, as there is no loss of information. [Currently, this may not
work everywhere, e.g. formal to actual parameters.] Narrowing, on the other hand, requires the
programmer to use explicit type conversion. In the following example, the variable one, a 32 bit
variable, is assigned to the variable two which has a size of 16 bits.. The following example shows how
the variable is properly converted using the name of the conversion type. In this case, the name _uint16
which is that of the variable two.

```
var one:   _uint32;
var two:   _uint16;
one = two;                      // implicitly widened
two = _uint16(one);             // explicitly truncated
```

5.2.2.2 Reference Coercion

A reference coercion is a automatic conversion from a base type, **T**, to a reference type **@T**. Since
ESL has no "address of" operator, reference coercions allow the "address of" to take place only at
assignment points. An example:

```
proc foo(arg: @[]_byte);
var string: [8]_byte;
. . .
foo(string);            // convert type []_byte to @[]_byte
```

5.3 Fundamental Terms

Precedence as presented in the following sections....

Operator	Description	Associativity
.	Package dereference	left-to-right
() _abs _man _min	Parentheses (function call) Absolute value Maximum value Minimum value	left-to-right
. . [] @ ?	Record dereference Method introduction Brackets (array subscript) Pointer dereference Property inquiry	left-to-right
+ - ~ !	Unary plus (no-operation) Unary minus Bitwise negation Boolean not	right-to-left
* / % << >> &	Multiplication Division Modulus Bitwise shift left Bitwise shift right Bitwise AND	left-to-right
+ - | ^	Addition Subtraction Bitwise OR Bitwise XOR	left-to-right
== != < <= > >=	Relational equal to Relational not equal to Relational less than Relational less than or equal to Relational greater than Relational greater than or equal to	left-to-right
&&	Boolean AND	left-to-right
||	Boolean OR	left-to-right

5.3.1 Type Queries

An identifier which has been declared as a type or variable may be queried to get attributes relating to its type. The result is a compile-time constant.

```
min
max
size
bits
align
len – applies only to arrays
```

An example:

```
type a: _uint8;
var b: 0..31;
var n: _uint;
type x: [10]_int;

a = a?max;              // value of 'a' is 255
a = a?min;              // value of 'a' is 0
n = a?size;             // value of 'n' is 1
n = a?bits;             // value of 'n' is 8
b = b?max;              // value of 'b' is 31
b = b?min;              // value of 'b' is 0
n = b?size;             // value of 'n' is 1
n = b?bits;             // value of 'n' is 5
n = x?len;              // value of 'n' is 10
```

5.4 Unary Operations

negation, logical inversion, boolean not
- (negation) – only applicable to signed types
~ (one's complement) – only applicable to unsigned types
! (boolean not) – only applicable to boolean type

5.5 Binary Operations

5.5.1 Multiplicative Operations

multiply, divide, modulo, shifts, and

```
*
/
%
<<
>>
&
```

5.5.2 Additive Operations

add, sub, xor, or

```
+
-
^
|
```

5.5.3 Comparison Operations

All comparison operations result in the built-in boolean type. Reference types and array types support only equal and not-equal operations. For array types, depending on length and alignment, a call to a library routine may be generated. Record types cannot be compared.

```
==
!=
<
<=
>
>=
```

5.5.4 Boolean And

```
&&
```

5.5.5 Boolean Or

```
||
```

5.6 Intrinsic Operations

There are several built-in operations which use a function call syntax.

5.6.1 Defined

The _defined operator returns true if the symbol is defined.

5.6.2 Absolute Value

The _abs operator takes the absolute value of an signed integral number. For example:

```
var a: _int32;
var b: _uint32;

a = -1234;
b = _abs(a);            // value of 'b' is 1123
```

5.6.3 Maximum Value

The _max operator takes the largest value of two integral numbers. For example:

```
var a: _int32;
var b: _int32;
var c: _int32;

a = 1234;
b = 5678;
c = _max(a, b);        // value of 'c' is 5678
```

5.6.4 Minimum Value

The _min operator takes the smallest value of two integral numbers. For example:

```
var a: _int32;
var b: _int32;
var c: _int32;

a = 1234;
b = 5678;
c = _min(a, b);        // value of 'c' is 1234
```

5.6.5 Count Leading Zeros

The _clz and _clznz operators calculate the number of leading zeros in the expression. When it is known that the expression is not equal to zero, the operator _clznz may be more efficient.

5.6.6 Count Trailing Zeros

The _ctz and _ctznz operators calculate the number of trailing zeros in the expression. When it is known that the expression is not equal to zero, the operator _ctznz may be more efficient.

5.6.7 Population Count

The operator _pop calculates the number of one bits in the expression.

5.6.8 Zero

The procedure _zero zeroes the contents of an instantiation of a type. See the following section for an example.

5.6.9 New Allocation

The _new operator allocates memory for a new instantiation of a type. For types with an unknown array length, the length of the "unknown" part must be specified in parentheses after the type name. An optional second parameter allows for an implementation specific option (e.g., choice of address space or memory pool selection). The operator _new returns memory with unspecified contents. A library call will be generated for each instance of _new.

Examples:

```
type rec:
{   a:      _uint;
    b:      [4]_uint;
};
type ident:
{   length: _uint;
    name:   []_byte;
};
var pr: @rec;
var pi: @ident;
pr = _new(rec);
_zero(pr);                      // zero new allocation
pi = _new(ident(8), 1);         // use of length and optional parameter
_zero(pi(8));                   // zero with length
```

5.6.10   Delete Allocation

The _delete operator de-allocates memory for a previously allocated type. As with _new, an "unknown" length must be explicitly specified. This allows for memory allocation algorithms which do not store the size of allocated memory in meta-data. Again, as with _new, there is an optional second parameter with the same meaning.

Examples, based on the examples in the previous section:

```
_delete(pe);
_delete(pr1(16));
_delete(pr2(8), 1);
```

6   Statements

All statements, except for assignment statements, begin with a keyword. All statements end with a semicolon.

Statement groups start with a "{" and end with a "}". There is no semicolon after a statement group.

6.1 Declarative Statements

6.1.1 Type Statement

The type statement associates a name with type. Refer to the previous section on types (section 3) for examples.

6.1.2 Variable Statement

The var statement instantiates an instance of a type. See the previous section on variables (section 4) for examples.

6.1.3 Alias Statement

The alias statement can be used for two distinct purposes. The most common use is to simply

introduce a new name into the current namespace which is a short-cut to another name. This is usually used to simplify access to names in a package.

The other use of an alias statement is to fix up a forward reference made in a package to some identifier outside that package (perhaps in a subsequent package).

Some examples are:

```
package foo
{ type R:
  { next: @R;           // self referential
    what: @P;           // forward reference
  };
}

package bar
{ alias foo.R as R;    // simplify
  type P:
  { next: @P;           // self referential
    from: @R;           // points to R in package foo
  };
  // fixup the forward reference in package foo
  alias P as foo.P;
}
```

6.2 Assignment Statement

6.2.1 Scalar Assignment

The most common example of scalar assignment is when a variable, this case a, is assigned a value that is a constant or the value of another variable.

```
a = 0;
a = b;
```

ESL also allows the assignment of multiple variables in a single statement.

```
a, b = 0, 0;
c, d = a, b;
```

In this example, the first statement assigns the variables a and b a value of zero, while the second statement assigns the values of c and d with the values of a and b. After the execution of the second statement the values of c and d would be zero, respectively.

As will be discussed in more detail in Section 6: Procedures, a procedure can return multiple values. The following example shows how the variables a and b are assigned the values that are returned from the procedure sumdiff.

```
a, b = sumdiff();
```

A more complex example shows a combination of assignments from variables and from the return values of the procedure sumdiff.

```
b, x, y, a = a, sumdiff(a, b), b;
```

In all the examples in this section, we have seen one or more variables assigned values with a corresponding number of variables on the right hand side of the assignment statement. Either from a variable or procedure or a combination of the two. In all cases, there has been a one-for-one assignment of a variable from another variable or from the return of a procedure. ESL does not allow for the splitting of a variable into composite types for assignment. In other words, two 16-bit variables cannot be assigned a value from a 32-bit variable on the right hand side. Conversely, a 32-bit variable cannot be assigned the values of two 16-bit variables on the right hand side.

6.2.2 Increment and Decrement Operators

Operations to increment or decrement a variable are so common they have their own syntax. Examples:

```
i += 1;                 // increment by 1
x -= 10;                // decrement by 10
```

6.2.3 Record Copy

Record assignment by copy is handled by the assignment statement:

```
type Record: { a: _uint; b: boolean; c: _byte; };
var r1, r2: Record;
. . .
r1 = r2;                // record copy
```

6.2.4 Array Copy

Array assignment by copy is handled by the assignment statement. This is especially powerful when array slices are involved:

```
type Array: [256]_byte;
var a1, a2: Array;
var i, j: 0..255;
. . .
a1 = a2;                // copy entire array
a1[0:8] = a2[16:8];     // copy 8 bytes from index 16 to index 0
a1[i:4] = a2[j:4];      // copy 4 bytes from index j to index i
```

6.3 Control Statements

6.3.1 If Statement

There are two variants of if statements, the boolean if and the selection if. Both variants have an optional else clause. The following is and example of a boolean if:

```
if x != 0 || x > y then
    y += x;
else
    y = 0;
```

The selection if allows for multiple conditions. For example:

```
if c
is 'A'..'Z' then
    x = c − 'A';
is 'a'..'z' then
    x = c − 'a';
is '_', ' ', '0'..'9' then
    x = 0;
else
    x = 0xFF;
```

6.3.2 While Statement

The while statement is a straightforward loop with the test at the top.

```
while x != 0 do
{
    . . .
}
```

6.3.3 Do Statement

The do statement is a straightforward loop with the test at the bottom.

```
do
{
    . . .
} while x != y;
```

6.3.4 For Statement

There are two variations of the for statement. The general form implies an order to evaluation. An example is:

```
for i from 0 to n-1 by 2 do
    x[i] += x[i+1];
```

The by clause is optional and defaults to "by 1".

The "type" form of the for statement specifies the range by either an explicit range clause or by use of a type name. This form allows the implementation to evaluate in any order. Two examples follow:

```
    for i in 0..n-1 do x[i] = 0;

    type Enum: (A, B, C, D, E, F);
    var array: [Enum] boolean;
    var ev: Enum;
    for ev in Enum do array[ev] = false;
```

6.3.5 Loop Statement

The loop statement creates a (possibly infinite) loop, unless there are exit statements within the loop:

```
    loop
    { . . .
      exit i == 0;
        i -= 1;
    }
```

6.3.6 Exit Statement

Exit statements must be enclosed in either a loop, while, or do statement. When the boolean expression evaluates to true, the loop is exited. Additional code may optionally be executed as the exit is taken.

```
    exit x <= 0;            // simple exit
    exit x <= y with        // exit which executes code when exit is taken
    {   x = -1;
        y = 0;
    }
```

6.3.7 Return Statement

Return statements evaluate returned values (if any) and then cause a procedure return.

```
    return;                 // return in a procedure with no return value
    return x-1;             // return in a procedure with one return value
    return a-b, b-a;        // return in a procedure with two return values
```

6.4 Asm Statement

The asm startment is used to insert an assembly language instruction.

The first string is the assembly language prototype. The second (optional) string is the constraint list. The statement ends with a list of expressions to be used as arguments. Variables in the prototype are indicated by "$" followed by a number, starting with zero.

```
    asm "halt";                          // no arguments
    asm "fire $0", "r", arg;             // one input register argument
    // the following has four arguments, two register inputs, and
    // two register outputs
    asm "sumdiff $0 $1 $2 $3", "=r,=r,r,r", ret1, ret2, arg1, arg2;
```

6.5 Error Statement

An error statement causes a compile-time error. This is most often used in conjunction with conditional compilation (see section 9.2).

7   Procedures

Procedures may return zero or more values.

7.1 Normal Procedures

Normal procedures are as expected. In the following example, the procedure 'sumdif' is defined to have two parameters and return two values:

```
proc sumdiff(a:_int, b:_int):_int, _int
{
    return a+b, a-b;
}

var x, y: _int;
...
x, y = sumdiff(x, y);
```

7.2 Methods

Methods allow the programmer to write a procedure associated with a particular type. (This is similar to "methods" in other languages, but the binding is to a type not a "class".) In the following example, the type 'foo' is declared and a method that operates on that type is defined:

```
type foo: 0..100;

proc (a: foo) saturate(b: _uint): foo
{   var sum: _uint;
    sum = a + b;
    sum = _min(sum, foo?max);
    return foo(sum);
}

var x: foo;
...
x = x.saturate(10);
```

While the previous example showed method on a scalar type, a more usual case is when the method type is a reference to a record:

```
type rec: { a:_uint; b:_boolean };

proc (p:@rec) init(x:_uint)
{
    p.a = x;
    p.b = true;
}

var r: rec;
...
r.init(0);                 // r will be reference coerced to @rec
```

7.3 Procedure Attributes

7.3.1 Inline

The *inline* attribute is currently only a hint to the compiler to tell it to "try very hard to inline" this
procedure whereever it is called. Procedures without the *inline* attribute may also be inlined, unless a
compiler flag is used to disable inlining.

7.3.2 External

The *external* attribute for procedures is used in a similar manner as it is for variables (see section
4.2.1).

```
package foo
{
    proc bar(): : external;
    proc baz(): _int: external("foobar");
    proc getcpuid(result: @[]_uint32): : external(0x1fff1ff1);
};
```

Here the code for 'getcpuid' is known to be located at address 0x1fff1ff1.

7.3.3 Global

The *global* attribute for procedures is used in a similar manner as it is for variables (see section 4.2.2.

7.3.4 Segment

[Not yet implemented.]

7.4 Forward Procedures

In some cases, a procedure must be called before it is defined (as with mutually recursive procedures).
A procedure may be declared forward of its use but before its definition by omitting the body of the
procedure and replacing it with a single semicolon:

```
proc sumdiff(a:_int, b:_int):_int, _int;
```

7.5 Procedure References

A procedure reference is a data type which contains a pointer to a procedure with a given "signature". The signature is the parameter list together with the returned value list. This is prefixed by '@_'.

```
proc sumdiff(a:_int, b:_int):_int, _int;
. . .
type pref: @_(a:_int, b:_int):_int, _int;
var p: pref;
p = sumdiff;
```

References to methods have the same syntax as those to normal procedures. The first parameter is that of the method parameter.

```
type bpref: @_(a:@rec, x:_uint);
var bp: bpref;
p = r.init;
```

8   Packages

Packages are the primary method of reusing code by providing a private namespace. A package is a bundle of declarations and procedures wrapped in a package name. Packages have no semantics other than providing an encapsulating namespace.

Examples:

```
package foo:
{ type type1: 0..1000;
  type type2: (no, yes, maybe);
  proc sub1(x: type1, y:type 1): type2
  { statement-list }
}
```

8.1 Package Continuation

Packages can be re-opened to extend the contents of the package. This is known as *package continuation*. An example is:

```
package foo
{
    var x: _uint32;
    var y: _uint32;
}

// code outside package foo

package foo              // continue package foo
{
    var z: _uint32;
}           // package foo now has variables x,y,z
```

A package can be imported into another package. This is not *package continuation*, but rather the definition of a new package containing a nested package.

Supposed a file named "pkg_A.esl" contains:

```
package pkg_A
{
    var v1: _uint32;
    var v2: _uint32;
};
```

In a different file, is the following code:

```
package pkg_B
{
    import pkg_A;
    var v1: _uint32;
    var v2: _uint16;
    proc x()
    {
        v1 = 33;
        v2 = 44;
        pkg_A.v1 = 55;
        pkg_B.v2 = 66:
    }
};
```

9   Programs and Scope

Programs consist of global declarations, packages, procedures and import statements.

9.1 Import Statement

The import statement causes file inclusion. The included file typically contains one or more packages, but this is not a requirement. The import statement can take an immediate symbol or an expression, within parentheses, that yields a compile-time constant which is a string.

```
import foo;           // includes file foo.esl
const bar = "foobar";
import (bar);         // includes file foobar.esl
```

Package specifiers can be long. When a package is imported, the package name can be renamed with the `alias` statement. This enables short names to be used in the program.

Alternatively, individual identifiers in the package can the aliased. In that way, the package name does not have to be used as a prefix on every reference to an identifier within that package.

9.2 Conditional Compilation

The if statement, as defined in section 6.3.1, can be used outside of procedures to implement conditional compilation. For example:

```
if _defined(debug) then {
   type foo: debug;
   proc ShowAll(arg: debug) { . . . }
}
```

Another example uses the error statement to trigger a compile-time error:

```
if _uint?size
is 8 then type foo: _uint64;
is 4 then type foo: _uint32;
else error "Size of _uint not handled";
```

Statements or statement groups that are not selected by the if statement are skipped over. This skipping takes place on a lexical token basis, so that everything that is skipped must consist of legal lexical tokens.

Any compile-time constant can be used as a selector expression, an advantage over those languages that use a separate pre-processor.

9.3 Compilation Style

Most programmers have adopted an incremental compile approach to program development. In other words, one compiles numerous source files, one at a time, and then links the resultant object files together to form a single, executable file. In the ESL / LLVM ecosystem, this approach is not necessary nor desirable.

Program development with ESL is best done with a 'compile all at once' approach. Simply put, one file imports all the remaining files in the program. This approach to compilation gives the ESL front-end, and to a greater extent the LLVM back-end, a complete view of all the software in a development project. This allows ESL and the LLVM tools the ability to maximize the optimization of the code generation, which can lead to advantages in code size and execution efficiency:

1. Procedures (non-global) that are not called do not generate any code.

2. Variables and constants (non-global) are not allocated.

3. Opportunities to in-line are maximized.

4. Constant propagation across procedure boundaries can lead to elimination of unneeded code.

The following shows how five files in a development are combined and compiled all at once, using ESL. The main file in this development is: top_level.esl. It imports the four other files in the development into top_level.esl.

=== top_level.esl

```
import file1;
import file2;
import file3;
import file4;
```

```
proc main(argc: _uint, argv: @[]@[] _byte): _int
{
  . . .
}
```

10 Language Syntax Summary

10.1   Notation

The syntax is specified in EBNF (Extended Backus-Naur Form). Lexical symbols are enclosed in double quotes.

10.2   Program Syntax

program　　　= *prog-stmt* { *prog-stmt*}
prog-stmt　　= *package* | *procedure* |
　　　　　　　import-stmt | *decl_stmt* | *if-stmt* | *error-stmt*
import-stmt　= "`import`" *file-specifier* "`;`"
error-stmt　　= "`error`" *string* "`;`"

10.3   Package Syntax

package　　　= "`package`" *ident* "`{`" *pkg-stmt* { *pkg-stmt* } "`}`"
pkg_stmt　　 = *package* | *procedure* |
　　　　　　　import-stmt | *decl_stmt* | *if-stmt* | *error-stmt*

10.4   Procedure Syntax

procedure　　= "`proc`" ["`(`" *formal* "`)`"] *ident* *signature* ["`:`" *proc-attr*]
　　　　　　　("`;`" | *group*)
signature　　 = "`(`" [*parm-list*] "`)`" ["`:`" [*retv-list*]]
parm-list　　 = *formal* { "`,`" *formal* }
formal　　　 = *ident* "`:`" *type-def*
retv-list　　 = *type-def* { `,` *type-def* }
group　　　　= "`{`" *stmt* { *stmt* } "`}`"
stmt　　　　 = *group* |
　　　　　　　decl_stmt |
　　　　　　　asgn-stmt | *if-stmt* |
　　　　　　　loop-stmt | *while-stmt* | *do-stmt* | *for-stmt* |
　　　　　　　exit-stmt | *return-stmt* | *error-stmt*

10.5 Declaration Statement Syntax

```
decl-stmt       = alias-stmt | type-stmt | var-stmt | const-stmt
alias-stmt      = "alias" alias-list ";"
alias-list      = alias-spec { "," alias-spec }
alias-spec      = qname "as" ( ident | pkg-name "." name )
type-stmt       = "type" ident [ "(" typeid ")" ] ":" type-def [ ":" type-attr-list ] ";"
var-stmt        = "var" ident { "," ident } ":" type-def [ ":" var-attr-list ] ";"
const-stmt      = "const" ident [ ":" type-def ] "=" constant ";"
type-def        = type-name | type-range | type-enum | type-ref | type-record | type-array
type-range      = number ".." number
type-enum       = "(" enum-list ")"
enum-list       = enum-def { "," enum-def }
enum-def        = ident [ "=" sconst ]
type-ref        = "@" type-def
type-record     = "{" field { field } "}"
field           = ident ":" type-def [ ":" attr-list ] ";"
type-array      = "[" type-index "]" type-def
type-name       = name
type-attr-list  = type-attr { , type-attr }
type-attr       = "bits" "(" cexpr ")" | "size" "(" cexpr ")" | "align" "(" cexpr ")" |
                  "lsb" | "msb" | "le" | "be" | "ro" | "wo" | "in" | "out"
var-attr        = "global" [ "(" x-attr ")" ] | "external" [ "(" x-attr ")" ]
x-attr          = string | cexpr
proc-attr       = "inline" | var-attr
```

10.6 Executable Statement Syntax

```
asgn-stmt       = lhs-list "=" expr-list ";"
                  lhs "+=" expr ";" | lhs "-=" expr ";"
if-stmt         = "if" bool-expr "then" stmt [ "else" stmt |
                  "if" expr is-list [ "else" stmt ]
is-list         = is-clause [ is-list ]
is-clause       = "is" is-value-list "then" stmt
is-value-list   = is-value { "," is-value-list }
is-value        = cexpr [ ".." cexpr ]
loop-stmt       = "loop" stmt
while-stmt      = "while" bool-expr "do" stmt
do-stmt         = "do" stmt "while" bool-expr
for-stmt        = "for" var-name for-suffix
for-suffix      = for-in | for-from
for-in          = "in" for-type "do" stmt
for-type        = type-name | subrange
for-from        = "from" expr "to" expr [ "by" expr ] "do" stmt
exit-stmt       = "exit" bool-expr [ "then" stmt ]
return-stmt     = "return" expr-list ";"
```

asm-stmt	= "asm" *string* ["," *string* { "," *expr* }] ";"
error-stmt	= "error" *string* ";"

10.7 Expression Syntax Summary

cexpr	= *expr* # compile time constant
expr	= *baexpr* { "||" *baexpr* }
baexpr	= *cmpexpr* { "&&" *cmpexpr* }
cmpexpr	= *addexpr* [*cmpop addexpr*]
cmpop	= "=="|"!="|"<"|">"|"<="|">="
addexpr	= *mulexpr* { *addop addexpr* }
addop	= "+"|"-"|"|"|"^"
mulexpr	= *uexpr* { *mulop mulexpr* }
mulop	= "*"|"/"|"%"|"<<"|">>"|"&"
uexpr	= *term* | *unop term*
unop	= "+"|"-"|"~"|"!"
term	= type-query | *number* | *qname* | "(" *expr* ")" | *cast*
type-query	= [*type-name* | *var-name*] "?" *type-qattr*
type-qattr	= "min"|"max"|"bits"|"size"|"align"|"len"
cast	= *type-name* "(" *expr* ["," *selector*] ")"

11 Library Support

```
memcpy
memcmpN
memalloc
memfree
```

12 ESL For C Programmers

Since C has been the dominant system programming language for the past 30 years or so, this chapter is designed to help C programmers to understand how to express C idioms as ESL idioms. For those programmers who have been exposed to "Pascal-like" languages (e.g. Module, Ada, etc.), some of this might be familiar.

One of the major differences between C and ESL is that ESL has no reserved words! You are free to name a variable "if" or a type "then" or a procedure "else".

12.1 Preprocessor

ESL does not have a separate pre-processor language. There are no features such as "token-pasteing."

12.1.1 Include

The ESL "package" and "import" features replace the C "header files" and"#include" mechanism.

12.1.2   If and ifdef

There is no equivalent to "#ifdef". However the "#if" pre-processor feature has an ESL equivalent. The if statement can be used if its selection expression is a compile-time constant. The ESL if statement deals in units of complete statements, not tokens.

12.2   Declarations

A major difference between C and ESL is that ESL uses the Pascal-like syntax for declarations. In ESL the identifer comes first followed by the type. In C, the reverse is true.

12.3   Typedefs

The ESL type statement is the equivalent of the C typedef.

12.4   Enumerations

The identifiers representing values in a C enumerations are exposed in the same name space as the enumeration. In ESL (as in some other languages) an enumeration constant identifier is private to the enumeration and exposed by giving its type name as a prefix. However, in many cases, the type can be inferred and the enumeration prefix is not necessary.

12.5   Pointers and Arrays

C example:
```
   typedef int *pi;                 // pointer to int
   typedef int ai[8];               // array of int
   typedef int *api[8];             // array of pointers to int
   typedef int (*pai)[8];           // pointer to array of int
```
ESL example:
```
   type pi: @_int;                  // pointer to int
   type ai: [8]_int;                // array of int
   type api: [8]@_int;              // array of pointers to int
   type pai: @[8]_int;              // pointer to array of int
```

12.6   Pointers to Procedures

In C, the declaration of a pointer to a function has a syntax that has confused even experienced programmers. Consider a pointer to a function that takes two integer arguments and returns an integer:
```
   typedef int (*func)(int, int);
```
In ESL, the syntax is still a little funky due to the use of the null identifier "_", but perhaps easier to read:
```
   type func: @_(a:_int, b:_int): _int;
```
Or the example from page 122 of the 2[nd] edition of Kernighan and Ritchie, x is definied as an "array[3] of pointer to function returning pointer to array[5] of char":
```
   typedef char (*(*x[3])())[5];
```

Which is ESL is written pretty much as the text describes it:

```
type x: [3]@_(): @[5]_byte;
```

12.7   Parameters and Arrays

In C function parameters are passed "by value" *except for arrays*. In the case of an array, the address of the array is passed instead. In ESL all procedure parameters are passed "by value", with no exceptions. If you want a pointer to an array, then the procedure parameter must be so declared.

Also, in C, pointer arithmetic may be used as a mechanism to access elements in arrays. This is not allowed in ESL, access to the element must be done with array indexing.

For example, the Unix command line convention passes an a pointer to an array of pointers to strings as an argument vector (e.g. "argv").   In C this is often specified as:

```
int main(int argc, char *argv[])
{
}
```

In ESL all the initial pointers and the specification of a string as "[]_byte" must be explicit:

```
proc main(argc: _int, argv: @[]@[]_byte): _int
{
}
```

12.8   Statements

12.8.1   Assignment Statements

Multiple assignment is not allowed in ESL. C statements such as
```
x = y = 0;
```
must be split into multiple assignment statements in ESL as in the following:

```
x = 0;
y = 0;
```

Alternatively, as mentioned in page 21, the following could be used:

```
x, y = 0, 0;
```

12.8.2   Assignment Within an Expression

In ESL, expressions cannot have side effects (except for procedure calls), such C constructs must be written as a separate assignment statement.

12.8.3   If Statements

Boolean if statements in ESL are very similar to those in C. The syntax is a little different, the parentheses surrounding the expression are not necessary, but the then keyword is.

12.8.4 Switch Statements

ESL doesn't have an explicit statement type corresponding to the C "switch" statement. The ESL "if" statement covers the same territory with one important difference - ESL doesn't allow "fall-through". Also each ESL "case" is a single statement, if more than one is required than a statement group delimited by "{" and "}" is
required.

C Example:

```
unsigned int foo;
switch (foo) {
  case 0: case 1: case 7: case 8: case 9: case 15:
    statement1;
    break;
  case 2: case 3: case 4:
    statement2;
    statement3;
  case 5:        /* fall thru */
    statement4;
    break;
  default:
    statement5;
    break;
}
```

ESL Example:

```
var foo: 0..15;
if foo
is 0..1, 7..9, 15 then
  statement1;
is 2..4 then
{ statement2;
  statement3;
  statement4;  // no fall thru
}
is 5 then
  statement4;
else
  statement5;
```

13 Building and Using the ESL Compiler

13.1 Prerequisites

Before one builds the ESL compiler, two packages must be install: the LLVM Compiler Infrastructure

and the GNU Compiler Collection (only the assembler and libc are used) for the target platform. The building and/or installation of both of these packages are described on their respective web pages and will not be repeated here.

13.2　Building the ESL compiler

13.2.1　Getting the source

After the LLVM and GNU collections have been installed, the source to ESL can be downloaded from the Google Project Hosting site: http://code.google.com/p/esl. The user has two alternative ways to get the source to ESL: either as a tarball or from the svn depository. The tarballs are available under the 'Downloads' tab, while instructions for checking out the svn despository are under the 'Source' tab.

13.2.2　Installing LLVM

13.2.3　Building the compiler

Once the sources to ESL have been obtained, the next step is to actually build and install the compiler. The top level directory of the distribution should contain two files: LICENSE and README, and two directories: src and doc. The building and installing of the compiler will occur in the src directory and all of the following directions will assume that the user has that directory as the current working directory.

Before starting the build process, a line must be changed in the shell script file: llvm. The line that set the shell variable LLVM must be changed to the directory where the LLVM tools, i.e., llc, llvm-as, etc., were installed into. The line looks like the following:

      LLVM=$HOME/work/llvm/Release+Asserts/bin   # !!FIX THIS!!

Once this line has been changed, building can proceed.

After the line in LLVM has been changed, the make(1) variables: CFLAGS, ARCH, and GCC should be examined and possibly adjusted for the target platform. The file: Makefile is the makefile used to build the ESL compiler and contains the definitions of the aforementioned variables. Once these two files have been examined and optionally changed, the ESL compiler can be built.

To build the ESL compiler, the user invokes 'make'. The make'ing of the ESL will proceed and produce four (4) versions of the compile: eslc0, eslc1, eslc2, and eslc3. The user should select the 'eslc3' version as the ESL compiler to be used for compilation. This binary can be put into the user's choice of a directory and making sure the directory defined in the user's shell path variable.

13.3　ESL Command Line Options

```
eslc  [-mtarget] [-Idir] [-ofile] [-D[asf]] [-M] [-g] files
```

13.3.1 -D[asftm]

Control debug output:

    `-a` Dump the abstract syntax tree
    `-s Dump the symbol table`
    `-f` List all source files as they are imported

13.3.2 -m*target*

Specify the target architecture.

ESL can only support what LLVM has implemented as target backends.

Current targets

x86

x86_64

x86_64-darwin

msp430

cortex-m3

arm920t

ppc32

ppc64

s390x

systemz

mips

13.3.3 -I*dir*

Add a directory to the set of directories search for import files. There may be multiple instances, they are searched in the order specified. For example:

    `-I.. -I../.. -Ilib`

causes the search path for imports to look first in the current directory, then in the parent directory, then in the parent of the parent directory and finally in the directory lib in the current directory.

13.3.4 -o*file*

13.3.5 -M

Instead of compiling, just output a list of files suitable for makefile dependencies.

13.3.6   -g

Generate debugging information for use by run-time debuggers such as gdb. [This doesn't work yet.]

13.4   Compilation Flow

eslc – the compiler "front-end"

llvm-as – converts llvm textual format to llvm byte-code

opt – architecture independent optimization

llc – the compiler "back-end" that generates assembly language [Future versions of LLVM may generate object code and eliminate the need for a target assembler.

gas – target assemble

ld – linker used to link in library routine